

# Blur Test

1.0.0

Generated by Doxygen 1.9.6



<b>1 Blur-Test</b>	<b>1</b>
<b>2 File Index</b>	<b>3</b>
2.1 File List	3
<b>3 File Documentation</b>	<b>5</b>
3.1 blur.cu File Reference	5
3.1.1 Detailed Description	7
3.1.2 Macro Definition Documentation	7
3.1.2.1 PBSTR	7
3.1.2.2 PBWIDTH	7
3.1.2.3 SAFE_CALL	7
3.1.3 Function Documentation	8
3.1.3.1 add_value() [1/2]	8
3.1.3.2 add_value() [2/2]	8
3.1.3.3 call_gaussian_blur_1d()	9
3.1.3.4 call_gaussian_blur_2d()	9
3.1.3.5 gaussian_blur() [1/2]	9
3.1.3.6 gaussian_blur() [2/2]	10
3.1.3.7 gaussian_blur_exit()	10
3.1.3.8 gaussian_blur_init()	11
3.1.3.9 gaussian_blur_x()	11
3.1.3.10 gaussian_blur_y()	12
3.1.3.11 generate_gaussian_kernel_1d()	12
3.1.3.12 generate_gaussian_kernel_2d()	13
3.1.3.13 main()	13
3.1.3.14 multiply_value() [1/3]	13
3.1.3.15 multiply_value() [2/3]	14
3.1.3.16 multiply_value() [3/3]	14
3.1.3.17 printProgress()	15
3.1.3.18 set_value() [1/5]	15
3.1.3.19 set_value() [2/5]	15
3.1.3.20 set_value() [3/5]	16
3.1.3.21 set_value() [4/5]	16
3.1.3.22 set_value() [5/5]	16
3.1.3.23 stress_test()	17
3.1.3.24 subtract_value() [1/2]	17
3.1.3.25 subtract_value() [2/2]	17
3.1.4 Variable Documentation	18
3.1.4.1 ginput	18
3.1.4.2 goutput	18
3.2 blur.cu	18
3.3 main.cpp File Reference	23

3.3.1 Detailed Description . . . . .	24
3.3.2 Macro Definition Documentation . . . . .	24
3.3.2.1 PBSTR . . . . .	24
3.3.2.2 PBWIDTH . . . . .	24
3.3.3 Function Documentation . . . . .	24
3.3.3.1 apply_convolution() . . . . .	24
3.3.3.2 apply_convolution_multi_threaded() . . . . .	25
3.3.3.3 apply_kernel() . . . . .	25
3.3.3.4 apply_kernel_multithreaded() . . . . .	26
3.3.3.5 generate_gaussian_kernel() . . . . .	26
3.3.3.6 main() . . . . .	26
3.3.3.7 printProgress() . . . . .	27
3.3.3.8 stress_test() . . . . .	27
3.4 main.cpp . . . . .	27
<b>Index</b>	<b>31</b>

## Chapter 1

# Blur-Test

Testing out some blurs with opecv, OpenMP and CUDA



## Chapter 2

# File Index

### 2.1 File List

Here is a list of all documented files with brief descriptions:

<a href="#">blur.cu</a>	Convolution blurring in Nvidia CUDA . . . . .	5
<a href="#">main.cpp</a>	Gaussian blurring using CPU . . . . .	23





## Chapter 3

# File Documentation

### 3.1 blur.cu File Reference

convolution blurring in Nvidia CUDA

```
#include <cstring>
#include <cuda_profiler_api.h>
#include <cuda_runtime.h>
#include <iostream>
#include <opencv2/core/core.hpp>
#include <opencv2/core/cuda.hpp>
#include <opencv2/core/cuda/common.hpp>
#include <opencv2/core/matx.hpp>
#include <opencv2/cudaimgproc.hpp>
#include <opencv2/highgui.hpp>
#include <opencv2/opencv.hpp>
#include <stdio.h>
```

#### Macros

- `#define PBSTR "||||||||||||||||||||||||||||||||||||||||||||||||||||||||"`
- `#define PBWIDTH 60`
- `#define SAFE_CALL(call, msg) _safe_cuda_call((call), (msg), __FILE__, __LINE__)`  
*a macro for sage calling CUDA functions*

#### Functions

- void `printProgress` (double percentage)
- `__host__` void `generate_gaussian_kernel_2d` (float \*kernel, const int n, const float sigma=1)  
*generate the gaussian kernel with given kernel size and standard deviation*
- `__host__` void `generate_gaussian_kernel_1d` (float \*kernel, const int n, const float sigma=1)  
*generate a 1D gaussian kernel*
- `__device__` `__forceinline__` void `set_value` (const int &val, uchar &out)  
*Sets the value of a uchar type.*
- `__device__` `__forceinline__` void `set_value` (const float &val, float &out)

- set the value for a floating point type.*

  - `__device__ __forceinline__ void set\_value (const float &val, float3 &out)`  
*set the value for a float3 tupe. All the 3 fields will have the value `val`*
  - `__device__ __forceinline__ void set\_value (const float3 &val, uchar3 &out)`  
*set the value for a unsigned char3 type with a float3 type*
  - `__device__ __forceinline__ void set\_value (const int &val, uchar3 &out)`  
*Sets the value of a uchar3 type.*
  - `__device__ __forceinline__ uchar3 subtract\_value (uchar3 in1, uchar3 in2)`  
*Subtraction for uchar3 types.*
  - `__device__ __forceinline__ float3 add\_value (float3 in1, float3 in2)`  
*add two values and return it*
  - `__device__ __forceinline__ float add\_value (float in1, float in2)`  
*add two floating point values*
  - `__device__ __forceinline__ uchar subtract\_value (uchar in1, uchar in2)`  
*Subtraction for uchar types.*
  - `__device__ __forceinline__ float3 multiply\_value (const float &x, const uchar3 &y)`  
*multiplication for float and uchar3 types. Multiply each filed in uchar3 with the float value and return a float3*
  - `__device__ __forceinline__ float3 multiply\_value (const float &x, const float3 &y)`  
*multiplication for float and float3 types. Multiply each filed in uchar3 with the float value and return a float3*
  - `__device__ __forceinline__ float multiply\_value (const float &x, const uchar &y)`  
*multiplication for float and uchar4 types*
- `template<typename T_in , typename T_out , typename F_cal >  
__global__ void gaussian\_blur (const float *kernel, int n, const cv::cuda::PtrStepSz< T_in > input, cv::cuda::PtrStepSz< T_out > output)`  
*applies the gaussian blur convolution to the input image*
- `template<typename T_in , typename T_out , typename F_cal >  
__global__ void gaussian\_blur\_x (float *kernel, int kernel_size, const cv::cuda::PtrStepSz< T_in > input, cv::cuda::PtrStepSz< T_out > output)`  
*applies the gaussian blur convolution to the input image along the x-axis*
- `template<typename T_in , typename T_out , typename F_cal >  
__global__ void gaussian\_blur\_y (float *kernel, int kernel_size, const cv::cuda::PtrStepSz< T_in > input, cv::cuda::PtrStepSz< T_out > output)`  
*applies the gaussian blur convolution to the input image along the y-axis*
- `template<typename... Ts>  
void gaussian\_blur\_exit (bool remove_globals, Ts &&...inputs)`  
*free all the GPU resources*
- `void call\_gaussian\_blur\_2d (float *d_kernel, const int &n, const cv::cuda::GpuMat &input, cv::cuda::GpuMat &output)`  
*calls the gaussian\_blur function appropriately based on the type of image*
- `void call\_gaussian\_blur\_1d (float *d_kernel, const int &n, const cv::cuda::GpuMat &input, cv::cuda::GpuMat &output)`  
*calls the separable gaussian\_blur function appropriately based on the type of image*
- `__host__ void gaussian\_blur (const cv::Mat &input, cv::Mat &output, const int n=3, const float sigma=1.0, bool two_d=true, bool remove_globals=true)`  
*the gaussian blur function which runs on the HOST CPU. It calls the `call_gaussian_blur` function after initialization of the appropriate values and kernel.*
- `void gaussian\_blur\_init (const cv::Mat &input, cv::Mat &output)`  
*initialization for gaussian blurring operation*
- `void stress\_test (const int &n, const bool &two_d)`
- `int main (int argc, char **argv)`

## Variables

- `cv::cuda::GpuMat` [ginput](#)
- `cv::cuda::GpuMat` [goutput](#)

### 3.1.1 Detailed Description

convolution blurring in Nvidia CUDA

#### Author

Arjun31415

Definition in file [blur.cu](#).

### 3.1.2 Macro Definition Documentation

#### 3.1.2.1 PBSTR

```
#define PBSTR "||||||||||||||||||||||||||||||||||||||||"
```

Definition at line 23 of file [blur.cu](#).

#### 3.1.2.2 PBWIDTH

```
#define PBWIDTH 60
```

Definition at line 24 of file [blur.cu](#).

#### 3.1.2.3 SAFE\_CALL

```
#define SAFE_CALL(  
    call,  
    msg ) _safe_cuda_call((call), (msg), __FILE__, __LINE__)
```

a macro for sage calling CUDA functions

#### Parameters

<i>call</i>	the CUDA function call
<i>msg</i>	user specified message

Definition at line 60 of file [blur.cu](#).

### 3.1.3 Function Documentation

#### 3.1.3.1 `add_value()` [1/2]

```
__device__ __forceinline__ float add_value (  
    float in1,  
    float in2 )
```

add two floating point values

##### Parameters

<i>in1</i>	value 1
<i>in2</i>	value 2

##### Returns

the sum

Definition at line 220 of file [blur.cu](#).

#### 3.1.3.2 `add_value()` [2/2]

```
__device__ __forceinline__ float3 add_value (  
    float3 in1,  
    float3 in2 )
```

add two values and return it

##### Parameters

<i>in1</i>	input 1
<i>in2</i>	input 2

##### Returns

returns the added value

Definition at line 208 of file [blur.cu](#).

### 3.1.3.3 call\_gaussian\_blur\_1d()

```
void call_gaussian_blur_1d (
    float * d_kernel,
    const int & n,
    const cv::cuda::GpuMat & input,
    cv::cuda::GpuMat & output )
```

calls the separable gaussian\_blur function appropriately based on the type of image

#### Parameters

<i>d_kernel</i>	the kernel, stored on GPU device memory
<i>n</i>	the size of the kernel
<i>input</i>	the input image stored on the GPU
<i>output</i>	the output image stored on the GPU

Definition at line 466 of file [blur.cu](#).

### 3.1.3.4 call\_gaussian\_blur\_2d()

```
void call_gaussian_blur_2d (
    float * d_kernel,
    const int & n,
    const cv::cuda::GpuMat & input,
    cv::cuda::GpuMat & output )
```

calls the gaussian\_blur function appropriately based on the type of image

#### Parameters

<i>d_kernel</i>	the kernel, stored on GPU device memory
<i>n</i>	the size of the kernel
<i>input</i>	the input image stored on the GPU
<i>output</i>	the output image stored on the GPU

Definition at line 433 of file [blur.cu](#).

### 3.1.3.5 gaussian\_blur() [1/2]

```
__host__ void gaussian_blur (
    const cv::Mat & input,
    cv::Mat & output,
    const int n = 3,
    const float sigma = 1.0,
```

```
bool two_d = true,
bool remove_globals = true )
```

the gaussian blur function which runs on the HOST CPU. It calls the `call_gaussian_blur` function after initialization of the appropriate values and kernel.

#### Parameters

<i>input</i>	the input image stored on the CPU memory
<i>output</i>	the output image stored on the CPU memory
<i>n</i>	the size of the Gaussian kernel, defaults to 3
<i>sigma</i>	the standard deviation of the Gaussian kernel, defaults to 1.
<i>two↔ _d</i>	whether to use the 2D gaussian blur kernel or two separable 1D gaussian blur kernels, defaults to true

Definition at line 506 of file [blur.cu](#).

#### 3.1.3.6 gaussian\_blur() [2/2]

```
template<typename T_in , typename T_out , typename F_cal >
__global__ void gaussian_blur (
    const float * kernel,
    int n,
    const cv::cuda::PtrStepSz< T_in > input,
    cv::cuda::PtrStepSz< T_out > output )
```

applies the gaussian blur convolution to the input image

#### Template Parameters

<i>t_in</i>	the type of input image, i.e uchar for black and white, uchar3 for rgb, float3 etc
<i>t_out</i>	the type of output image
<i>f_cal</i>	the type for calculating intermediate sums and products

#### Parameters

<i>kernel</i>	the kernel to apply the convolution
<i>n</i>	the dimension of the kernel ( $n \times n$ )
<i>input</i>	the input image
<i>output</i>	the output image

Definition at line 290 of file [blur.cu](#).

#### 3.1.3.7 gaussian\_blur\_exit()

```
template<typename... Ts>
void gaussian_blur_exit (
```

```
bool remove_globals,
Ts &&... inputs )
```

free all the GPU resources

#### Template Parameters

<i>Ts</i>	
-----------	--

#### Parameters

<i>inputs</i>	varidaic list of resources
<i>remove_globals</i>	if true, removes the global variables, otherwise not, if Not then user has to handle the removal of the global variables and freeing the GPU memory

Definition at line 414 of file [blur.cu](#).

### 3.1.3.8 gaussian\_blur\_init()

```
void gaussian_blur_init (
    const cv::Mat & input,
    cv::Mat & output )
```

initialization for gaussian blurring operation

#### Parameters

<i>input</i>	input image stored on the CPU
<i>output</i>	output image stored on the CPU

Definition at line 549 of file [blur.cu](#).

### 3.1.3.9 gaussian\_blur\_x()

```
template<typename T_in , typename T_out , typename F_cal >
__global__ void gaussian_blur_x (
    float * kernel,
    int kernel_size,
    const cv::cuda::PtrStepSz< T_in > input,
    cv::cuda::PtrStepSz< T_out > output )
```

applies the gaussian blur convolution to the input image along the x-axis

#### Template Parameters

<i>t_in</i>	the type of input image, i.e uchar for black and white, uchar3 for rgb, float3 etc
<i>t_out</i>	the type of output image
<i>F_cal</i>	the type for calculating intermediate sums and products

## Parameters

<i>kernel</i>	the kernel to apply the convolution
<i>kernel_size</i>	the dimension of the kernel
<i>input</i>	the input image
<i>output</i>	the output image

Definition at line 336 of file [blur.cu](#).

### 3.1.3.10 gaussian\_blur\_y()

```
template<typename T_in , typename T_out , typename F_cal >
__global__ void gaussian_blur_y (
    float * kernel,
    int kernel_size,
    const cv::cuda::PtrStepSz< T_in > input,
    cv::cuda::PtrStepSz< T_out > output )
```

applies the gaussian blur convolution to the input image along the y-axis

## Template Parameters

<i>t_in</i>	the type of input image, i.e uchar for black and white, uchar3 for rgb, float3 etc
<i>t_out</i>	the type of output image
<i>f_cal</i>	the type for calculating intermediate sums and products

## Parameters

<i>kernel</i>	the kernel to apply the convolution
<i>kernel_size</i>	the dimension of the kernel
<i>input</i>	the input image
<i>output</i>	the output image

Definition at line 377 of file [blur.cu](#).

### 3.1.3.11 generate\_gaussian\_kernel\_1d()

```
__host__ void generate_gaussian_kernel_1d (
    float * kernel,
    const int n,
    const float sigma = 1 )
```

generate a 1D gaussian kernel



## Parameters

<i>kernel</i>	the array in which the weights are stored
<i>n</i>	the size of the kernel. a 1D kernel of length n is needed
<i>sigma</i>	the standard deviation of the kernel

## Returns

Definition at line 106 of file [blur.cu](#).

**3.1.3.12 generate\_gaussian\_kernel\_2d()**

```
__host__ void generate_gaussian_kernel_2d (
    float * kernel,
    const int n,
    const float sigma = 1 )
```

generate the gaussian kernel with given kernel size and standard deviation

## Parameters

<i>kernel</i>	the array in which the weights are stored
<i>n</i>	the size of the kernel, t.e. n x n kernel is needed
<i>sigma</i>	the standard deviation

Definition at line 70 of file [blur.cu](#).

**3.1.3.13 main()**

```
int main (
    int argc,
    char ** argv )
```

Definition at line 571 of file [blur.cu](#).

**3.1.3.14 multiply\_value() [1/3]**

```
__device__ __forceinline__ float3 multiply_value (
    const float & x,
    const float3 & y )
```

multiplication for float and float3 types. Multiply each filed in uchar3 with the float value and return a float3

**Parameters**

$x$	Input 1
$y$	Input 2

**Returns**

value after multiplication

Definition at line 259 of file [blur.cu](#).

**3.1.3.15 multiply\_value() [2/3]**

```
__device__ __forceinline__ float multiply_value (
    const float & x,
    const uchar & y )
```

multiplication for float and uchar4 types

**Parameters**

$x$	Input 1
$y$	Input 2

**Returns**

$x*y$

Definition at line 272 of file [blur.cu](#).

**3.1.3.16 multiply\_value() [3/3]**

```
__device__ __forceinline__ float3 multiply_value (
    const float & x,
    const uchar3 & y )
```

multiplication for float and uchar3 types. Multiply each filed in uchar3 with the float value and return a flolat3

**Parameters**

$x$	Input 1
$y$	Input 2

**Returns**

value after multiplication

Definition at line 245 of file [blur.cu](#).

**3.1.3.17 printProgress()**

```
void printProgress (
    double percentage )
```

Definition at line 25 of file [blur.cu](#).

**3.1.3.18 set\_value() [1/5]**

```
__device__ __forceinline__ void set_value (
    const float & val,
    float & out )
```

set the value for a floating point type.

**Parameters**

<i>val</i>	the value
<i>out</i>	the output

Definition at line 142 of file [blur.cu](#).

**3.1.3.19 set\_value() [2/5]**

```
__device__ __forceinline__ void set_value (
    const float & val,
    float3 & out )
```

set the value for a float3 tupe. All the 3 fields will have the value *val*

**Parameters**

<i>val</i>	the value
<i>out</i>	the output

Definition at line 154 of file [blur.cu](#).

**3.1.3.20 set\_value()** [3/5]

```
__device__ __forceinline__ void set_value (
    const float3 & val,
    uchar3 & out )
```

set the value for a unsigned char3 type with a flot3 type

**Parameters**

<i>val</i>	the value to set
<i>out</i>	the ouput

Definition at line [165](#) of file [blur.cu](#).

**3.1.3.21 set\_value()** [4/5]

```
__device__ __forceinline__ void set_value (
    const int & val,
    uchar & out )
```

Sets the value of a uchar type.

**Parameters**

<i>val</i>	The value
<i>out</i>	The output

Definition at line [131](#) of file [blur.cu](#).

**3.1.3.22 set\_value()** [5/5]

```
__device__ __forceinline__ void set_value (
    const int & val,
    uchar3 & out )
```

Sets the value of a uchar3 type.

**Parameters**

in	<i>val</i>	The value
	<i>out</i>	The output

Definition at line [177](#) of file [blur.cu](#).

### 3.1.3.23 stress\_test()

```
void stress_test (
    const int & n,
    const bool & two_d )
```

Definition at line 554 of file [blur.cu](#).

### 3.1.3.24 subtract\_value() [1/2]

```
__device__ __forceinline__ uchar subtract_value (
    uchar in1,
    uchar in2 )
```

Subtraction for uchar types.

#### Parameters

in	<i>in1</i>	Input 1
in	<i>in2</i>	Input 2

#### Returns

Output

Definition at line 233 of file [blur.cu](#).

### 3.1.3.25 subtract\_value() [2/2]

```
__device__ __forceinline__ uchar3 subtract_value (
    uchar3 in1,
    uchar3 in2 )
```

Subtraction for uchar3 types.

#### Parameters

in	<i>in1</i>	Input 1
in	<i>in2</i>	Input 2

#### Returns

Output

Definition at line 192 of file [blur.cu](#).

### 3.1.4 Variable Documentation

#### 3.1.4.1 ginput

cv::cuda::GpuMat ginput

Definition at line 20 of file [blur.cu](#).

#### 3.1.4.2 goutput

cv::cuda::GpuMat goutput

Definition at line 20 of file [blur.cu](#).

## 3.2 blur.cu

[Go to the documentation of this file.](#)

```

00001
00007 #include <cstring>
00008 #undef __noinline__
00009 #include <cuda_profiler_api.h>
00010 #include <cuda_runtime.h>
00011 #include <iostream>
00012 #include <opencv2/core/core.hpp>
00013 #include <opencv2/core/cuda.hpp>
00014 #include <opencv2/core/cuda/common.hpp>
00015 #include <opencv2/core/matx.hpp>
00016 #include <opencv2/cudaimgproc.hpp>
00017 #include <opencv2/highgui.hpp>
00018 #include <opencv2/opencv.hpp>
00019 #include <stdio.h>
00020 cv::cuda::GpuMat ginput, goutput;
00021
00022 // Progress Bar STRing
00023 #define PBSTR "|||||
00024 #define PBWIDTH 60
00025 void printProgress(double percentage)
00026 {
00027     int val = (int)(percentage * 100);
00028     int lpad = (int)(percentage * PBWIDTH);
00029     int rpad = PBWIDTH - lpad;
00030     printf("\r%3d% [%.*s%s]", val, lpad, PBSTR, rpad, "");
00031     fflush(stdout);
00032 }
00033
00043 static inline void _safe_cuda_call(cudaError err, const char *msg,
00044                                     const char *file_name, const int line_number)
00045 {
00046     if (err != cudaSuccess)
00047     {
00048         fprintf(stderr, "%s\n\nFile: %s\n\nLine Number: %d\n\nReason: %s\n",
00049             msg, file_name, line_number, cudaGetErrorString(err));
00050         std::cin.get();
00051         exit(EXIT_FAILURE);
00052     }
00053 }
00060 #define SAFE_CALL(call, msg) _safe_cuda_call((call), (msg), __FILE__, __LINE__)
00061
00070 __host__ void generate_gaussian_kernel_2d(float *kernel, const int n,
00071                                           const float sigma = 1)
00072 {
00073     int mean = n / 2;
00074     float sumOfWeights = 0;

```

```

00075     float p, q = 2.0 * sigma * sigma;
00076
00077     // Compute weights
00078     for (int i = 0; i < n; i++)
00079     {
00080         for (int j = 0; j < n; j++)
00081         {
00082             p = sqrt((i - mean) * (i - mean) + (j - mean) * (j - mean));
00083             kernel[i * n + j] = std::exp(-(p * p) / q) / (M_PI * q);
00084             sumOfWeights += kernel[i * n + j];
00085         }
00086     }
00087
00088     // Normalizing weights
00089     for (int i = 0; i < n; i++)
00090     {
00091         for (int j = 0; j < n; j++)
00092         {
00093             kernel[i * n + j] /= sumOfWeights;
00094         }
00095     }
00096 }
00097
00106 __host__ void generate_gaussian_kernel_1d(float *kernel, const int n,
00107                                           const float sigma = 1)
00108 {
00109     // Calculate the values of the kernel
00110     float sum = 0.0f;
00111     for (int i = 0; i < n; i++)
00112     {
00113         float x = i - (n - 1) / 2.0f;
00114         kernel[i] = std::exp(-x * x / (2 * sigma * sigma));
00115         sum += kernel[i];
00116     }
00117
00118     // Normalize the kernel so that its sum equals 1
00119     for (int i = 0; i < n; i++)
00120     {
00121         kernel[i] /= sum;
00122     }
00123 }
00124
00131 __device__ __forceinline__ void set_value(const int &val, uchar &out)
00132 {
00133     out = val;
00134 }
00135
00142 __device__ __forceinline__ void set_value(const float &val, float &out)
00143 {
00144     out = val;
00145 };
00146
00154 __device__ __forceinline__ void set_value(const float &val, float3 &out)
00155 {
00156     out.x = val, out.y = val, out.z = val;
00157 }
00158
00165 __device__ __forceinline__ void set_value(const float3 &val, uchar3 &out)
00166 {
00167     out.x = val.x;
00168     out.y = val.y;
00169     out.z = val.z;
00170 }
00177 __device__ __forceinline__ void set_value(const int &val, uchar3 &out)
00178 {
00179     out.x = val;
00180     out.y = val;
00181     out.z = val;
00182 }
00183
00192 __device__ __forceinline__ uchar3 subtract_value(uchar3 in1, uchar3 in2)
00193 {
00194     uchar3 out;
00195     out.x = in1.x - in2.x;
00196     out.y = in1.y - in2.y;
00197     out.z = in1.z - in2.z;
00198     return out;
00199 }
00200
00208 __device__ __forceinline__ float3 add_value(float3 in1, float3 in2)
00209 {
00210     return {in1.x + in2.x, in1.y + in2.y, in1.z + in2.z};
00211 }
00212
00220 __device__ __forceinline__ float add_value(float in1, float in2)
00221 {
00222     return in1 + in2;

```

```

00223 }
00224
00233 __device__ __forceinline__ uchar subtract_value(uchar in1, uchar in2)
00234 {
00235     return in1 - in2;
00236 }
00245 __device__ __forceinline__ float3 multiply_value(const float &x,
00246                                                    const uchar3 &y)
00247 {
00248     return {x * (float)y.x, x * (float)y.y, x * (float)y.z};
00249 }
00250
00259 __device__ __forceinline__ float3 multiply_value(const float &x,
00260                                                    const float3 &y)
00261 {
00262     return {x * (float)y.x, x * (float)y.y, x * (float)y.z};
00263 }
00264
00272 __device__ __forceinline__ float multiply_value(const float &x, const uchar &y)
00273 {
00274     return x * (float)y;
00275 }
00276
00289 template <typename T_in, typename T_out, typename F_cal>
00290 __global__ void gaussian_blur(const float *kernel, int n,
00291                               const cv::cuda::PtrStepSz<T_in> input,
00292                               cv::cuda::PtrStepSz<T_out> output)
00293 {
00294     // calculate the x & y position of the current image pixel
00295     const int x = blockIdx.x * blockDim.x + threadIdx.x;
00296     const int y = blockIdx.y * blockDim.y + threadIdx.y;
00297
00298     if (x >= input.cols || y >= input.rows) return;
00299
00300     const int mid = n / 2;
00301     F_cal sum;
00302     set_value(0, sum);
00303     // synchronize all the threads till this point
00304     __syncthreads();
00305
00306     // loop over the n x n neighborhood of the current pixel
00307     for (int i = 0; i < n; i++)
00308     {
00309         for (int j = 0; j < n; j++)
00310         {
00311             int y_idx = y + i - mid;
00312             int x_idx = x + j - mid;
00313             if (y_idx > input.rows || x_idx > input.cols) continue;
00314             const float kernel_val = kernel[(n - i - 1) * n + (n - j - 1)];
00315             sum =
00316                 add_value(sum, multiply_value(kernel_val, input(y_idx, x_idx)));
00317         }
00318     }
00319     T_out result;
00320     set_value(sum, result);
00321     output(y, x) = result;
00322 }
00335 template <typename T_in, typename T_out, typename F_cal>
00336 __global__ void gaussian_blur_x(float *kernel, int kernel_size,
00337                                 const cv::cuda::PtrStepSz<T_in> input,
00338                                 cv::cuda::PtrStepSz<T_out> output)
00339 {
00340     int x = blockIdx.x * blockDim.x + threadIdx.x;
00341     int y = blockIdx.y * blockDim.y + threadIdx.y;
00342     const int radius = kernel_size / 2;
00343     const int width = input.cols;
00344     const int height = input.rows;
00345
00346     if (x >= input.cols || y >= input.rows) return;
00347
00348     F_cal pixel;
00349     set_value(0, pixel);
00350
00351     for (int i = -radius; i <= radius; i++)
00352     {
00353         int idx = y * width + (x + i);
00354         /* printf("%d\n", idx); */
00355         if (idx >= 0 && idx < width * height)
00356         {
00357             const float weight = kernel[i + radius];
00358             pixel = add_value(pixel, multiply_value(weight, input[idx]));
00359         }
00360     }
00361     set_value(pixel, output(y, x));
00362 }
00363
00376 template <typename T_in, typename T_out, typename F_cal>

```



```

00377 __global__ void gaussian_blur_y(float *kernel, int kernel_size,
00378                               const cv::cuda::PtrStepSz<T_in> input,
00379                               cv::cuda::PtrStepSz<T_out> output)
00380 {
00381     int x = blockIdx.x * blockDim.x + threadIdx.x;
00382     int y = blockIdx.y * blockDim.y + threadIdx.y;
00383     const int radius = kernel_size / 2;
00384     const int width = input.cols;
00385     const int height = input.rows;
00386
00387     if (x >= input.cols || y >= input.rows) return;
00388
00389     F_cal pixel;
00390     set_value(0, pixel);
00391     float weight_sum = 0;
00392     for (int i = -radius; i <= radius; i++)
00393     {
00394         int idx = (y + i) * width + x;
00395         if (idx >= 0 && idx < width * height)
00396         {
00397             float weight = kernel[i + radius];
00398             pixel = add_value(pixel, multiply_value(weight, input[idx]));
00399         }
00400     }
00401     set_value(pixel, output(y, x)); // output(y,x) = pixel;
00402 }
00403
00413 template <typename... Ts>
00414 void gaussian_blur_exit(bool remove_globals, Ts &&...inputs)
00415 {
00416     if (remove_globals)
00417     {
00418         ginput.release();
00419         goutput.release();
00420     }
00421     ([&] { SAFE_CALL(cudaFree(inputs), "Unable to free"); }(), ...);
00422 }
00423
00433 void call_gaussian_blur_2d(float *d_kernel, const int &n,
00434                           const cv::cuda::GpuMat &input,
00435                           cv::cuda::GpuMat &output)
00436 {
00437     CV_Assert(input.channels() == 1 || input.channels() == 3);
00438     const dim3 block(16, 16);
00439
00440     // Calculate grid size to cover the whole image
00441     const dim3 grid(cv::cuda::device::divUp(input.cols, block.x),
00442                   cv::cuda::device::divUp(input.rows, block.y));
00443     if (input.channels() == 1)
00444     {
00445         gaussian_blur<uchar, uchar, float>
00446             <<grid, block>>(d_kernel, n, input, output);
00447         return;
00448     }
00449     else if (input.channels() == 3)
00450     {
00451         gaussian_blur<uchar3, uchar3, float3>
00452             <<grid, block>>(d_kernel, n, input, output);
00453     }
00454     cudaSafeCall(cudaGetLastError());
00455 }
00466 void call_gaussian_blur_1d(float *d_kernel, const int &n,
00467                            const cv::cuda::GpuMat &input,
00468                            cv::cuda::GpuMat &output)
00469 {
00470     CV_Assert(input.channels() == 1 || input.channels() == 3);
00471     const int block_size = 16;
00472     dim3 dimBlock(block_size, block_size);
00473     dim3 dimGrid(cv::cuda::device::divUp(input.cols, dimBlock.x),
00474                 cv::cuda::device::divUp(input.rows, dimBlock.y));
00475     cv::cuda::GpuMat temp = input.clone();
00476     // Apply the horizontal Gaussian blur
00477     if (input.channels() == 1)
00478     {
00479         gaussian_blur_x<uchar, uchar, float>
00480             <<dimGrid, dimBlock>>(d_kernel, n, input, temp);
00481         gaussian_blur_y<uchar, uchar, float>
00482             <<dimGrid, dimBlock>>(d_kernel, n, temp, output);
00483     }
00484     else if (input.channels() == 3)
00485     {
00486         gaussian_blur_x<uchar3, uchar3, float3>
00487             <<dimGrid, dimBlock>>(d_kernel, n, input, temp);
00488         gaussian_blur_y<uchar3, uchar3, float3>
00489             <<dimGrid, dimBlock>>(d_kernel, n, temp, output);
00490     }
00491 }

```

```

00492     cudaSafeCall(cudaGetLastError());
00493 }
00506 __host__ void gaussian_blur(const cv::Mat &input, cv::Mat &output,
00507                             const int n = 3, const float sigma = 1.0,
00508                             bool two_d = true, bool remove_globals = true)
00509 {
00510     ginput.upload(input);
00511     std::vector<float> gauss_kernel_host;
00512     float *d_gauss_kernel;
00513     if (two_d)
00514     {
00515         gauss_kernel_host = std::vector<float>(n * n);
00516         generate_gaussian_kernel_2d(gauss_kernel_host.data(), n, sigma);
00517         cudaMalloc((void **)&d_gauss_kernel, n * n * sizeof(float));
00518         SAFE_CALL(cudaMemcpy(d_gauss_kernel, gauss_kernel_host.data(),
00519                             sizeof(float) * n * n, cudaMemcpyHostToDevice),
00520                 "Unable to copy kernel");
00521         /* cudaProfilerStart(); */
00522         call_gaussian_blur_2d(d_gauss_kernel, n, ginput, goutput);
00523         /* cudaProfilerStop(); */
00524     }
00525     else
00526     {
00527         gauss_kernel_host = std::vector<float>(n);
00528         generate_gaussian_kernel_1d(gauss_kernel_host.data(), n, sigma);
00529         cudaMalloc((void **)&d_gauss_kernel, n * sizeof(float));
00530         SAFE_CALL(cudaMemcpy(d_gauss_kernel, gauss_kernel_host.data(),
00531                             sizeof(float) * n, cudaMemcpyHostToDevice),
00532                 "Unable to copy kernel");
00533         /* cudaProfilerStart(); */
00534         call_gaussian_blur_1d(d_gauss_kernel, n, ginput, goutput);
00535         /* cudaProfilerStop(); */
00536     }
00537     // goutput.upload(input);
00538     goutput.download(output);
00539     gaussian_blur_exit(remove_globals, d_gauss_kernel);
00540 }
00541 }
00542
00549 void gaussian_blur_init(const cv::Mat &input, cv::Mat &output)
00550 {
00551     ginput.create(input.rows, input.cols, input.type());
00552     goutput.create(output.rows, output.cols, output.type());
00553 }
00554 void stress_test(const int &n, const bool &two_d)
00555 {
00556     std::cout << "Kernel size: " << n << std::endl;
00557     const std::string path = "../images/peppers_color.tif";
00558     cv::Mat input = cv::imread(path, 1);
00559     auto output = input.clone();
00560     gaussian_blur_init(input, output);
00561     for (int i = 0; i < 100; i++)
00562     {
00563         printProgress((float)i / 100);
00564         gaussian_blur(input, output, n, 1.7, two_d, false);
00565     }
00566     std::cout << std::endl;
00567     ginput.release();
00568     goutput.release();
00569     return;
00570 }
00571 int main(int argc, char **argv)
00572 {
00573     if (argc < 3)
00574     {
00575         printf("usage: Blur_Test <kernel_size> <Image_Path> [<Output_Path>]\n");
00576         return -1;
00577     }
00578     std::string mTitle = "Display Image";
00579     cv::Mat input;
00580     int n = atoi(argv[1]);
00581     if (strcmp(argv[2], "stress2d", 8) == 0)
00582     {
00583         stress_test(n, true);
00584         return 0;
00585     }
00586     else if (strcmp(argv[2], "stress1d", 8) == 0)
00587     {
00588         stress_test(n, false);
00589         return 0;
00590     }
00591     input = cv::imread(argv[2], 1);
00592     if (!input.data)
00593     {
00594         printf("No image data \n");
00595         return -1;
00596     }

```

```

00597     }
00598     auto output = input.clone();
00599
00600     // Call the wrapper function
00601     gaussian_blur_init(input, output);
00602     gaussian_blur(input, output, n, 1.7, 0);
00603
00604     // Show the input and output
00605     cv::imshow("Output", output);
00606
00607     // Wait for key press
00608     cv::waitKey();
00609     namedWindow(mTitle, cv::WINDOW_AUTOSIZE);
00610     /* namedWindow("gauss", cv::WINDOW_AUTOSIZE); */
00611     imshow(mTitle, input);
00612     /* imshow("gaussian", output); */
00613     if (argc >= 4) imwrite(argv[3], output);
00614     do
00615     {
00616
00617         auto k = cv::waitKey(500);
00618         if (k == 27)
00619         {
00620             cv::destroyAllWindows();
00621             return 0;
00622         }
00623         if (cv::getWindowProperty(mTitle, cv::WND_PROP_VISIBLE) == 0) return 0;
00624     } while (true);
00625     return 0;
00626
00627     return 0;
00628 }

```

## 3.3 main.cpp File Reference

gaussian blurring using CPU

```

#include <cmath>
#include <fstream>
#include <iostream>
#include <numeric>
#include <omp.h>
#include <opencv2/core.hpp>
#include <opencv2/highgui.hpp>
#include <opencv2/imgcodecs.hpp>
#include <opencv2/opencv.hpp>
#include <string>
#include <vector>

```

### Macros

- `#define PBSTR "||||||||||||||||||||||||||||||||||||||||||||||||||||||||"`
- `#define PBWIDTH 60`

### Functions

- void `printProgress` (double percentage)
- void `generate_gaussian_kernel` (std::vector< std::vector< float > > &kernel, const int n, const float sigma=1)  
Generate a 2D gaussian kernel.
- void `apply_convolution` (const std ::vector< std::vector< float > > &kernel, const Mat &original\_img, Mat &new\_img, const int &r, const int &c)  
apply a convolution kernel to a pixel

- void [apply\\_convolution\\_multi\\_threaded](#) (const std::vector< std::vector< float > > &kernel, const Mat &original\_img, Mat &new\_img, const int &r, const int &c)  
*apply a convolution kernel to a pixel using multiple threads (OMP)*
- void [apply\\_kernel](#) (const std::vector< std::vector< float > > &kernel, const Mat &original\_img, Mat &new\_img)  
*apply a convolution kernel to the entire image*
- void [apply\\_kernel\\_multithreaded](#) (const std::vector< std::vector< float > > &kernel, const Mat &original\_img, Mat &new\_img)  
*apply a convolution kernel to the entire image using multiple threads (OMP)*
- void [stress\\_test](#) (const int &n, const bool &multi=true)
- int [main](#) (int argc, char \*\*argv)

### 3.3.1 Detailed Description

gaussian blurring using CPU

Author

Arjun31415

Definition in file [main.cpp](#).

### 3.3.2 Macro Definition Documentation

#### 3.3.2.1 PBSTR

```
#define PBSTR "||||||||||||||||||||||||||||||||||||||||"
```

Definition at line 20 of file [main.cpp](#).

#### 3.3.2.2 PBWIDTH

```
#define PBWIDTH 60
```

Definition at line 21 of file [main.cpp](#).

### 3.3.3 Function Documentation

#### 3.3.3.1 apply\_convolution()

```
void apply_convolution (
    const std::vector< std::vector< float > > & kernel,
    const Mat & original_img,
    Mat & new_img,
    const int & r,
    const int & c )
```

apply a convolution kernel to a pixel

## Parameters

<i>kernel</i>	the convolution kernel
<i>original_img</i>	the original image
<i>new_img</i>	the output image
<i>r</i>	the row number of the current pixel
<i>c</i>	the column number of the current pixel

Definition at line 76 of file [main.cpp](#).

### 3.3.3.2 apply\_convolution\_multi\_threaded()

```
void apply_convolution_multi_threaded (
    const std::vector< std::vector< float > > & kernel,
    const Mat & original_img,
    Mat & new_img,
    const int & r,
    const int & c )
```

apply a convolution kernel to a pixel using multiple threads (OMP)

## Parameters

<i>kernel</i>	the convolution kernel
<i>original_img</i>	the original image
<i>new_img</i>	the output image
<i>r</i>	the row number of the current pixel
<i>c</i>	the column number of the current pixel

Definition at line 106 of file [main.cpp](#).

### 3.3.3.3 apply\_kernel()

```
void apply_kernel (
    const std::vector< std::vector< float > > & kernel,
    const Mat & original_img,
    Mat & new_img )
```

apply a convolution kernel to the entire image

## Parameters

<i>kernel</i>	the convolution kernel
<i>original_img</i>	the original image
<i>new_img</i>	the output image

Definition at line 136 of file [main.cpp](#).

#### 3.3.3.4 `apply_kernel_multithreaded()`

```
void apply_kernel_multithreaded (
    const std::vector< std::vector< float > > & kernel,
    const Mat & original_img,
    Mat & new_img )
```

apply a convolution kernel to the entire image using multiple threads (OMP)

##### Parameters

<i>kernel</i>	the convolution kernel
<i>original_img</i>	the original_img
<i>new_img</i>	the output image

Definition at line 155 of file [main.cpp](#).

#### 3.3.3.5 `generate_gaussian_kernel()`

```
void generate_gaussian_kernel (
    std::vector< std::vector< float > > & kernel,
    const int n,
    const float sigma = 1 )
```

Generate a 2D gaussian kernel.

##### Parameters

<i>kernel</i>	the kernel to be populated
<i>n</i>	the size of the kernel, the <code>kernel</code> must be of size $n * n$
<i>sigma</i>	the standard deviation of the gaussian kernel

Definition at line 39 of file [main.cpp](#).

#### 3.3.3.6 `main()`

```
int main (
    int argc,
    char ** argv )
```

Definition at line 241 of file [main.cpp](#).

### 3.3.3.7 printProgress()

```
void printProgress (
    double percentage )
```

Definition at line 23 of file [main.cpp](#).

### 3.3.3.8 stress\_test()

```
void stress_test (
    const int & n,
    const bool & multi = true )
```

Definition at line 170 of file [main.cpp](#).

## 3.4 main.cpp

[Go to the documentation of this file.](#)

```
00001
00007 #include <cmath>
00008 #include <fstream>
00009 #include <iostream>
00010 #include <numeric>
00011 #include <omp.h>
00012 #include <opencv2/core.hpp>
00013 #include <opencv2/highgui.hpp>
00014 #include <opencv2/imgcodecs.hpp>
00015 #include <opencv2/opencv.hpp>
00016 #include <string>
00017 #include <vector>
00018 using namespace cv;
00019
00020 #define PBSTR "|||||
00021 #define PBWIDTH 60
00022
00023 void printProgress(double percentage)
00024 {
00025     int val = (int)(percentage * 100);
00026     int lpad = (int)(percentage * PBWIDTH);
00027     int rpad = PBWIDTH - lpad;
00028     printf("\r%3d%% [%.*s%s]", val, lpad, PBSTR, rpad, "");
00029     fflush(stdout);
00030 }
00031
00039 void generate_gaussian_kernel(std::vector<std::vector<float>> &kernel,
00040                             const int n, const float sigma = 1)
00041 {
00042     int mean = n / 2;
00043     float sumOfWeights = 0;
00044     float p, q = 2.0 * sigma * sigma;
00045
00046     // Compute weights
00047     for (int i = 0; i < n; i++)
00048     {
00049         for (int j = 0; j < n; j++)
00050         {
00051             p = sqrt((i - mean) * (i - mean) + (j - mean) * (j - mean));
00052             kernel[i][j] = std::exp(-(p * p) / q) / (M_PI * q);
00053             sumOfWeights += kernel[i][j];
00054         }
00055     }
00056
00057     // Normalizing weights
00058     for (int i = 0; i < n; i++)
00059     {
00060         for (int j = 0; j < n; j++)
00061         {
00062             kernel[i][j] /= sumOfWeights;
```

```

00063     }
00064 }
00065 }
00076 void apply_convolution(const std::vector<std::vector<float>> &kernel,
00077                       const Mat &original_img, Mat &new_img, const int &r,
00078                       const int &c)
00079 {
00080     const size_t n = kernel.size();
00081     assert(n % 2 == 1);
00082     assert(n == kernel[0].size());
00083     const size_t mid = n / 2;
00084     new_img.at<Vec3b>(r, c) = {0, 0, 0};
00085     for (int i = 0; i < n; i++)
00086     {
00087         for (int j = 0; j < n; j++)
00088         {
00089             if (r - mid + i >= 0 && r - mid + i < original_img.rows &&
00090                 c - mid + j >= 0 && c - mid + j < original_img.cols)
00091                 new_img.at<Vec3b>(r, c) +=
00092                     kernel[n - i - 1][n - j - 1] *
00093                     original_img.at<Vec3b>(r - mid + i, c - mid + j);
00094         }
00095     }
00096 }
00106 void apply_convolution_multi_threaded(
00107     const std::vector<std::vector<float>> &kernel, const Mat &original_img,
00108     Mat &new_img, const int &r, const int &c)
00109 {
00110     const size_t n = kernel.size();
00111     assert(n % 2 == 1);
00112     assert(n == kernel[0].size());
00113     const size_t mid = n / 2;
00114     new_img.at<Vec3b>(r, c) = {0, 0, 0};
00115     #pragma omp parallel for shared(r, c, original_img, new_img, kernel)
00116     for (int i = 0; i < n; i++)
00117     {
00118         #pragma omp parallel for shared(r, c, original_img, new_img, kernel)
00119         for (int j = 0; j < n; j++)
00120         {
00121             if (r - mid + i >= 0 && r - mid + i < original_img.rows &&
00122                 c - mid + j >= 0 && c - mid + j < original_img.cols)
00123                 new_img.at<Vec3b>(r, c) +=
00124                     kernel[n - i - 1][n - j - 1] *
00125                     original_img.at<Vec3b>(r - mid + i, c - mid + j);
00126         }
00127     }
00128 }
00136 void apply_kernel(const std::vector<std::vector<float>> &kernel,
00137                  const Mat &original_img, Mat &new_img)
00138 {
00139     for (int i = 0; i < original_img.rows; i++)
00140     {
00141         for (int j = 0; j < original_img.cols; j++)
00142         {
00143             apply_convolution(kernel, original_img, new_img, i, j);
00144         }
00145     }
00146 }
00155 void apply_kernel_multithreaded(const std::vector<std::vector<float>> &kernel,
00156                                const Mat &original_img, Mat &new_img)
00157 {
00158     #pragma omp barrier
00159     #pragma omp parallel for shared(original_img, new_img, kernel)
00160     for (int i = 0; i < original_img.rows; i++)
00161     {
00162         #pragma omp parallel for shared(original_img, new_img, kernel)
00163         for (int j = 0; j < original_img.cols; j++)
00164         {
00165             apply_convolution(kernel, original_img, new_img, i, j);
00166         }
00167     }
00168     #pragma omp barrier
00169 }
00170 void stress_test(const int &n, const bool &multi = true)
00171 {
00172     std::cout << "Stress testing" << std::endl;
00173     const std::string path = "../images/peppers_color.tif";
00174     auto image = imread(path, 1);
00175     auto new_img = image.clone();
00176     std::vector<std::vector<float>> gauss_kernel(n, std::vector<float>(n));
00177     generate_gaussian_kernel(gauss_kernel, n, 1.6);
00178     std::vector<double> run_times;
00179     const int num_runs = 20;
00180     std::string fname;
00181     if (!multi)
00182     {
00183         fname = "profile_single_threaded.csv";

```



```

00184
00185     for (int i = 0; i < num_runs; i++)
00186     {
00187         printProgress((float)i / num_runs);
00188         auto start = std::chrono::high_resolution_clock::now();
00189         apply_kernel(gauss_kernel, image, new_img);
00190         auto end = std::chrono::high_resolution_clock::now();
00191         std::chrono::duration<double, std::milli> duration_ms = end - start;
00192         run_times.push_back(duration_ms.count());
00193     }
00194 }
00195 else
00196 {
00197     fname = "profile_multi_threaded.csv";
00198     for (int i = 0; i < num_runs; i++)
00199     {
00200         printProgress((float)i / num_runs);
00201         auto start = std::chrono::high_resolution_clock::now();
00202         apply_kernel_multithreaded(gauss_kernel, image, new_img);
00203         auto end = std::chrono::high_resolution_clock::now();
00204         std::chrono::duration<double, std::milli> duration_ms = end - start;
00205         run_times.push_back(duration_ms.count());
00206     }
00207 }
00208 /* for (auto i = 0; i < num_runs; i++)
00209     std::cout << run_times[i] << "\t"; */
00210 sort(run_times.begin(), run_times.end());
00211 double avg = std::accumulate(run_times.begin(), run_times.end(), 0.0) /
00212             run_times.size();
00213 double median =
00214     ((run_times.size() % 2 == 0) ? (run_times[run_times.size() / 2 - 1] +
00215                                     run_times[run_times.size() / 2]) /
00216      2
00217     : run_times[run_times.size() / 2]);
00218
00219 std::cout << "\nMax(ms)\t\tAvg(ms)\t\tMedian(ms)\t\tMin(ms)\n";
00220 std::cout << run_times.back() << "\t\t" << avg << "\t\t" << median
00221           << "\t\t" << run_times.front() << "\n";
00222 std::fstream file(fname, std::ios::in | std::ios_base::app);
00223 if (file.tellg() == 0)
00224 {
00225     // write the headers if the file is empty
00226     file << "KERNEL_SIZE,MAX_RUN_TIME,MIN_RUN_TIME,AVG_RUN_TIME,MEDIAN_RUN_"
00227           << "TIME"
00228           << std::endl;
00229 }
00230 file << n << "," << run_times.back() << "," << run_times.front() << ","
00231       << avg << "," << median << std::endl;
00232
00233 file.close();
00234
00235 /* setenv("MAX_RUN_TIME", std::to_string(run_times.back()).c_str(), 1);
00236 setenv("MIN_RUN_TIME", std::to_string(run_times.front()).c_str(), 1);
00237 setenv("AVG_RUN_TIME", std::to_string(avg).c_str(), 1);
00238 setenv("MEDIAN_RUN_TIME", std::to_string(median).c_str(), 1); */
00240 }
00241 int main(int argc, char **argv)
00242 {
00243     if (argc < 3)
00244     {
00245         printf("usage: Blur_Test <kernel_size> <Image_Path> [<Output_Path>]\n");
00246         return -1;
00247     }
00248     int n = atoi(argv[1]);
00249     if (strcmp(argv[2], "stressm", 7) == 0)
00250     {
00251         stress_test(n, true);
00252         return 0;
00253     }
00254     else if (strcmp(argv[2], "stress", 6) == 0)
00255     {
00256         stress_test(n, false);
00257         return 0;
00258     }
00259
00260     std::vector<std::vector<float>> gauss_kernel(n, std::vector<float>(n));
00261     generate_gaussian_kernel(gauss_kernel, n, 1.6);
00262
00263     std::string mTitle = "Display Image";
00264     Mat image;
00265     image = imread(argv[2], 1);
00266     if (!image.data)
00267     {
00268         printf("No image data \n");
00269         return -1;
00270     }

```

```
00271     namedWindow(mTitle, WINDOW_AUTOSIZE);
00272     auto new_img = image.clone();
00273     // namedWindow("gauss", WINDOW_AUTOSIZE);
00274     apply_kernel_multithreaded(gauss_kernel, image, new_img);
00275     imshow(mTitle, image);
00276     imshow("gaussian", new_img);
00277     if (argc >= 4) imwrite(argv[3], new_img);
00278     do
00279     {
00280
00281         auto k = waitKey(500);
00282         if (k == 27)
00283         {
00284             cv::destroyAllWindows();
00285             return 0;
00286         }
00287         if (cv::getWindowProperty(mTitle, WND_PROP_VISIBLE) == 0)
00288         {
00289             return 0;
00290             break;
00291         }
00292     } while (true);
00293     return 0;
00294 }
00295 }
```

# Index

- add\_value
  - blur.cu, [8](#)
- apply\_convolution
  - main.cpp, [24](#)
- apply\_convolution\_multi\_threaded
  - main.cpp, [25](#)
- apply\_kernel
  - main.cpp, [25](#)
- apply\_kernel\_multithreaded
  - main.cpp, [26](#)
- blur.cu, [5](#)
  - add\_value, [8](#)
  - call\_gaussian\_blur\_1d, [8](#)
  - call\_gaussian\_blur\_2d, [9](#)
  - gaussian\_blur, [9](#), [10](#)
  - gaussian\_blur\_exit, [10](#)
  - gaussian\_blur\_init, [11](#)
  - gaussian\_blur\_x, [11](#)
  - gaussian\_blur\_y, [12](#)
  - generate\_gaussian\_kernel\_1d, [12](#)
  - generate\_gaussian\_kernel\_2d, [13](#)
  - ginput, [18](#)
  - goutput, [18](#)
  - main, [13](#)
  - multiply\_value, [13](#), [14](#)
  - PBSTR, [7](#)
  - PBWIDTH, [7](#)
  - printProgress, [15](#)
  - SAFE\_CALL, [7](#)
  - set\_value, [15](#), [16](#)
  - stress\_test, [16](#)
  - subtract\_value, [17](#)
- call\_gaussian\_blur\_1d
  - blur.cu, [8](#)
- call\_gaussian\_blur\_2d
  - blur.cu, [9](#)
- gaussian\_blur
  - blur.cu, [9](#), [10](#)
- gaussian\_blur\_exit
  - blur.cu, [10](#)
- gaussian\_blur\_init
  - blur.cu, [11](#)
- gaussian\_blur\_x
  - blur.cu, [11](#)
- gaussian\_blur\_y
  - blur.cu, [12](#)
- generate\_gaussian\_kernel
  - main.cpp, [26](#)
- generate\_gaussian\_kernel\_1d
  - blur.cu, [12](#)
- generate\_gaussian\_kernel\_2d
  - blur.cu, [13](#)
- ginput
  - blur.cu, [18](#)
- goutput
  - blur.cu, [18](#)
- main
  - blur.cu, [13](#)
  - main.cpp, [26](#)
- main.cpp, [23](#)
  - apply\_convolution, [24](#)
  - apply\_convolution\_multi\_threaded, [25](#)
  - apply\_kernel, [25](#)
  - apply\_kernel\_multithreaded, [26](#)
  - generate\_gaussian\_kernel, [26](#)
  - main, [26](#)
  - PBSTR, [24](#)
  - PBWIDTH, [24](#)
  - printProgress, [26](#)
  - stress\_test, [27](#)
- multiply\_value
  - blur.cu, [13](#), [14](#)
- PBSTR
  - blur.cu, [7](#)
  - main.cpp, [24](#)
- PBWIDTH
  - blur.cu, [7](#)
  - main.cpp, [24](#)
- printProgress
  - blur.cu, [15](#)
  - main.cpp, [26](#)
- SAFE\_CALL
  - blur.cu, [7](#)
- set\_value
  - blur.cu, [15](#), [16](#)
- stress\_test
  - blur.cu, [16](#)
  - main.cpp, [27](#)
- subtract\_value
  - blur.cu, [17](#)